

Problem 1 (15 marks)

(Learning SVM via Co-ordinate Ascent) Consider the soft-margin linear SVM problem

$$\arg \max_{0 \leq \alpha \leq C} f(\alpha)$$

where $f(\alpha) = \alpha^\top \mathbf{1} - \frac{1}{2} \alpha^\top \mathbf{G} \alpha$, \mathbf{G} is an $N \times N$ matrix such that $G_{nm} = y_n y_m \mathbf{x}_n^\top \mathbf{x}_m$ and $\alpha = [\alpha_1, \alpha_2, \dots, \alpha_N]$ are the Lagrange multipliers. Given the optimal α , the SVM weight vector is $\mathbf{w} = \sum_{n=1}^N \alpha_n y_n \mathbf{x}_n$

Your goal is to derive a **co-ordinate ascent** procedure for the vector α , such that each iteration updates a uniformly randomly chosen entry α_n of the vector α . However, instead of updating α via standard co-ordinate descent as $\alpha_n = \alpha_n + \eta g_n$ where g_n denotes the n -th entry of the gradient vector $\nabla_\alpha f(\alpha)$, we will update it as $\alpha_n = \alpha_n + \delta_*$ where $\delta_* = \arg \max_\delta f(\alpha + \delta \mathbf{e}_n)$ and \mathbf{e}_n denotes a vector of all zeros except a 1 at entry n .

Essentially, this will give the new α_n that guarantees the maximum increase in f , with all other α_n 's fixed at their current value. Derive the expression for δ_* and give a sketch of the overall co-ordinate ascent algorithm.

Note that your expression for δ_* should be such that the constraint $0 \leq \alpha_n \leq C$ is maintained.

PS: I know I had said that, in this homework, I will give a programming task to implement SVM. :-) Well, even though I am not asking you to implement it, by solving the problem above, you would have pretty much everything you need to implement one of the state-of-the-art algorithms for linear SVM.

Problem 2 (5 marks)

(Within and Across) Suppose we wish to cluster some data by learning a function f such that $f_n = f(\mathbf{x}_n)$ is the cluster assignment for point \mathbf{x}_n . Show that finding f by minimizing \mathcal{L}_W , which is defined as the sum of squared distances between all pairs of points that are *within* the same cluster, i.e.,

$$\arg \min_f \mathcal{L}_W = \arg \min_f \sum_{n,m} \mathbb{I}[f_n = f_m] \|\mathbf{x}_n - \mathbf{x}_m\|^2$$

implicitly *also* maximizes the sum of squared distances between all pairs of points that are in *different* clusters. (Note: It is not for credit but you can also show that the above is equivalent to the K -means objective!)

Problem 3 (15 marks)

(Estimating a Gaussian when Data is Missing) Suppose we have collected N observations $\{\mathbf{x}_1, \dots, \mathbf{x}_N\}$ using a sensor. Let us assume each $\mathbf{x} \in \mathbb{R}^D$ as generated from a Gaussian distribution $\mathcal{N}(\mu, \Sigma)$. We would like to estimate the mean and covariance of this Gaussian. However, suppose the sensor was faulty and each \mathbf{x}_n could only have part of it as observed (think of a blacked out image). Denote $\mathbf{x}_n = [\mathbf{x}_n^{obs}, \mathbf{x}_n^{miss}]$ where \mathbf{x}_n^{obs} and \mathbf{x}_n^{miss} denote the observed and missing parts, respectively, of \mathbf{x}_n . We only get to see \mathbf{x}_n^{obs} . Note that different observations could have different parts as missing (e.g., different images may have different sets of pixels as missing), so the indices of the observed/missing entries of the vector \mathbf{x}_n may be different for different n .

Your goal is to develop an EM algorithm that gives maximum likelihood estimates of μ and Σ given this partially observed data. In particular, in the EM setting, you will treat each \mathbf{x}_n^{miss} as a latent variable and estimate its conditional distribution $p(\mathbf{x}_n^{miss} | \mathbf{x}_n^{obs}, \mu, \Sigma)$, given the current estimates μ and Σ of the parameters. In the M step, you will re-estimate μ and Σ , and will alternate between E and M steps until convergence.

Clearly write down the following: (1) The expression for $p(\mathbf{x}_n^{miss} | \mathbf{x}_n^{obs}, \mu, \Sigma)$; (2) The expected CLL for this model; (3) The update equations for μ and Σ .

Also clearly write down all the steps of the EM algorithm in this case, with appropriate equations.

Note/Hint: For this problem, you may find it useful to use the result that if $\mathbf{x} = [\mathbf{x}_a, \mathbf{x}_b]$ is Gaussian then $p(\mathbf{x}_a | \mathbf{x}_b)$ is also Gaussian (you may refer to Section 4.3.1 of MLAPP for the result).

Problem 4 (10 marks)

(Semi-supervised Classification) Consider learning a generative classification model for K -class classification with Gaussian class-conditionals $\mathcal{N}(\mathbf{x}|\mu_k, \Sigma_k)$, $k = 1, \dots, K$ with class marginals $p(y = k) = \pi_k$. However, unlike traditional generative classification, in this setting we are given N labeled examples $\{(\mathbf{x}_n, y_n)\}_{n=1}^N$ and an additional M *unlabeled* examples $\{\mathbf{x}_{N+1}, \dots, \mathbf{x}_{N+M}\}$. Design an EM algorithm to estimate all the unknowns of this model and clearly write down the expressions required in each step of the EM algorithm.

Note: You need not re-do the derivations we have done in the class or other homeworks/practice sets; feel free to re-use those without re-deriving from scratch.

Problem 5 (25 marks)

(Latent Variable Models for Supervised Learning) Consider learning a regression model given training data $\{(\mathbf{x}_n, y_n)\}_{n=1}^N$, with $\mathbf{x}_n \in \mathbb{R}^D$ and $y_n \in \mathbb{R}$. Let us give a small twist to the standard probabilistic linear model for regression that we have seen in the class. In particular, we will be introducing a latent variable z_n with each training example (\mathbf{x}_n, y_n) . The generative story would now be as follows

$$\begin{aligned} z_n &\sim \text{multinoulli}(\pi_1, \dots, \pi_K) \\ y_n &\sim \mathcal{N}(\mathbf{w}_{z_n}^\top \mathbf{x}_n, \beta^{-1}) \end{aligned}$$

Note that the model for the responses y_n is still discriminative, since the inputs are not being modeled.

The latent variables are $\mathbf{Z} = (z_1, \dots, z_N)$ and the global parameters are $\Theta = \{(\mathbf{w}_1, \dots, \mathbf{w}_K), (\pi_1, \dots, \pi_K)\}$.

(1) Give a brief explanation (max. 5 sentences) of what the above model is doing and why you might want to use it instead of the standard probabilistic linear model which models each response as $y_n \sim \mathcal{N}(\mathbf{w}^\top \mathbf{x}_n, \beta^{-1})$.

(2) Derive an ALT-OPT algorithm to estimate \mathbf{Z} and (MLE of) Θ , and clearly write down each step's update equations. For \mathbf{Z} , you must give the update equation for each individual latent variable $z_n, n = 1, \dots, N$. Likewise, for Θ , you must give the update equation for each $\mathbf{w}_k, k = 1, \dots, K$, and $\pi_k, k = 1, \dots, K$. Also, what will be the update of each z_n if $\pi_k = 1/K, \forall k$. Give a brief intuitive explanation (max 1-2 sentences) as to what this update does.

(3) Derive an expectation-maximization (EM) algorithm to estimate \mathbf{Z} and (MLE of) Θ , and clearly write down each step's update equations. Also show that, as $\beta \rightarrow \infty$, the EM algorithm reduces to ALT-OPT.

Note: It is okay to skip some of the standard/obvious steps from your derivations and write down the final expressions directly if the derivations are similar to what we have done in the class or previous homeworks/practice sets, but your final expressions must be clearly and unambiguously written.